

BME 646/ ECE695DL: Homework 3

Shaswata Roy

7 February 2022

1 Introduction

This assignment helps understand the use of various optimizers while performing stochastic gradient descent while updating the parameters.

2 Methodology

Computational graph primer was installed and after glancing through the code changes were made to optimize the stochastic gradient descent.

3 Implementation and Results

one_neuron_classifier_sgd_plus.py

```
import random
import numpy as np
import matplotlib.pyplot as plt
import operator
seed = 0
random.seed(seed)
np.random.seed(seed)
from ComputationalGraphPrimer import *

class ComputationalGraphPrimerSGDPlus(ComputationalGraphPrimer):
    def __init__(self, *args, **kwargs):
        if 'mu' in kwargs:
            self.mu = kwargs.pop('mu')
        else:
            self.mu = 0
        super().__init__(*args, **kwargs)
        self.prev_step=[]
```

```

def backprop_and_update_params_one_neuron_model(self, y_error,
vals_for_input_vars, deriv_sigmoid):
    input_vars = self.independent_vars
    vals_for_input_vars_dict = dict(zip(input_vars, list(vals_for_input_vars)))
    for i, param in enumerate(self.vals_for_learnable_params):
        step = self.mu*self.prev_step[i] + self.learning_rate * y_error
        *vals_for_input_vars_dict[input_vars[i]] * deriv_sigmoid
        self.vals_for_learnable_params[param] += step
        self.prev_step[i] = step
    self.step_size = self.mu*self.step_size
    +self.learning_rate*y_error*deriv_sigmoid
    self.bias += self.step_size

def reset_step(self):
    for i in range(len(self.learnable_params)):
        if len(self.prev_step) < len(self.vals_for_learnable_params):
            self.prev_step.append(0)
        else:
            self.prev_step[i] = 0

def run_training_loop_one_neuron_model(self, training_data):
    self.vals_for_learnable_params = {param: random.uniform(0,1)
    for param in self.learnable_params}
    self.bias = random.uniform(0,1)
    self.step_size = 0

class DataLoader:
    def __init__(self, training_data, batch_size):
        self.training_data = training_data
        self.batch_size = batch_size
        self.class_0_samples = [(item, 0) for item in self.training_data[0]]
        self.class_1_samples = [(item, 1) for item in self.training_data[1]]
    def __len__(self):
        return len(self.training_data[0]) + len(self.training_data[1])
    def _getitem(self):
        cointoss = random.choice([0,1])
        if cointoss == 0:
            return random.choice(self.class_0_samples)
        else:
            return random.choice(self.class_1_samples)
    def getbatch(self):
        batch_data, batch_labels = [], []
        maxval = 0.0
        for _ in range(self.batch_size):
            item = self._getitem()

```

```

        if np.max(item[0]) > maxval:
            maxval = np.max(item[0])
            batch_data.append(item[0])
            batch_labels.append(item[1])
        batch_data = [item/maxval for item in batch_data]
        batch = [batch_data, batch_labels]
        return batch

data_loader = DataLoader(training_data, batch_size=self.batch_size)
loss_running_record = []
i = 0
avg_loss_over_literations = 0.0
self.reset_step()
for i in range(self.training_iterations):
    data = data_loader.getbatch()
    data_tuples = data[0]
    class_labels = data[1]
    y_preds, deriv_sigmoids = self.forward_prop_one_neuron_model(data_tuples)
    loss = sum([(abs(class_labels[i] - y_preds[i]))**2
    for i in range(len(class_labels))])
    loss_avg = loss / float(len(class_labels))
    avg_loss_over_literations += loss_avg
    if i%(self.display_loss_how_often) == 0:
        avg_loss_over_literations /= self.display_loss_how_often
        loss_running_record.append(avg_loss_over_literations)
        print("[iter=%d] loss = %.4f" % (i+1, avg_loss_over_literations))
        avg_loss_over_literations = 0.0
    y_errors = list(map(operator.sub, class_labels, y_preds))
    y_error_avg = sum(y_errors) / float(len(class_labels))
    deriv_sigmoid_avg = sum(deriv_sigmoids) / float(len(class_labels))
    data_tuple_avg = [sum(x) for x in zip(*data_tuples)]
    data_tuple_avg = list(map(operator.truediv, data_tuple_avg,
        [float(len(class_labels))] * len(class_labels) ))
    self.backprop_and_update_params_one_neuron_model(y_error_avg,
        data_tuple_avg, deriv_sigmoid_avg)

return loss_running_record

if __name__ == '__main__':
    cgp = ComputationalGraphPrimerSGDPlus(
        one_neuron_model = True,
        expressions = ['xw=ab*xa+bc*xb+cd*xc+ac*xd'],
        output_vars = ['xw'],
        dataset_size = 5000,
        learning_rate = 1e-3,
        # learning_rate = 5 * 1e-2,

```

```

        training_iterations = 40000,
        batch_size = 8,
        display_loss_how_often = 100,
        debug = True
    )

    cgp.parse_expressions()
    training_data = cgp.gen_training_data()

    loss_sgd = cgp.run_training_loop_one_neuron_model( training_data )

    cgp_sgd_plus = ComputationalGraphPrimerSGDPlus(
        one_neuron_model = True,
        expressions = ['xw=ab*xa+bc*xb+cd*xc+ac*xd'],
        output_vars = ['xw'],
        dataset_size = 5000,
        learning_rate = 1e-3,
        #      learning_rate = 5 * 1e-2,
        training_iterations = 40000,
        batch_size = 8,
        display_loss_how_often = 100,
        debug = True,
        mu = 0.99
    )

    cgp_sgd_plus.parse_expressions()
    training_data = cgp_sgd_plus.gen_training_data()

    loss_sgd_plus = cgp_sgd_plus.run_training_loop_one_neuron_model( training_data )

    plt.figure()
    plt.plot(loss_sgd, label = "SGD Training Loss")
    plt.plot(loss_sgd_plus, label = "SGD+ Training Loss")
    plt.xlabel('iterations')
    plt.ylabel('loss')
    plt.title('SGD+ vs SGD Loss One Neuron')
    plt.legend()
    plt.show()

```

multi_neuron_classifier_sgd_plus.py

```

import random
import numpy as np
import matplotlib.pyplot as plt
import operator
seed = 0

```

```

random.seed(seed)
np.random.seed(seed)
from ComputationalGraphPrimer import *

class ComputationalGraphPrimerSGDPlus(ComputationalGraphPrimer):
    def __init__(self, *args, **kwargs):
        if 'mu' in kwargs:
            self.mu = kwargs.pop('mu')
        super().__init__(*args, **kwargs)
        self.prev_step=[]

    def backprop_and_update_params_multi_neuron_model(self, y_error, class_labels):
        # backproped prediction error:
        pred_err_backproped_at_layers = {i : [] for i in range(1,self.num_layers-1)}
        pred_err_backproped_at_layers[self.num_layers-1] = [y_error]
        for back_layer_index in reversed(range(1,self.num_layers)):
            input_vals = self.forw_prop_vals_at_layers[back_layer_index-1]
            input_vals_avg = [sum(x) for x in zip(*input_vals)]
            input_vals_avg = list(map(operator.truediv, input_vals_avg,
                [float(len(class_labels))] * len(class_labels)))
            deriv_sigmoid = self.gradient_vals_for_layers[back_layer_index]
            deriv_sigmoid_avg = [sum(x) for x in zip(*deriv_sigmoid)]
            deriv_sigmoid_avg = list(map(operator.truediv, deriv_sigmoid_avg,
                [float(len(class_labels))] * len(class_labels)))
            vars_in_layer = self.layer_vars[back_layer_index] ## a list of lists
            vars_in_next_layer_back = self.layer_vars[back_layer_index - 1] ## a list of lists

            layer_params = self.layer_params[back_layer_index]
            ## note that layer_params are stored in a dict like
            ##      {1: [['ap', 'aq', 'ar', 'as'], ['bp', 'bq', 'br', 'bs']], 2: [['cp',
            ## "layer_params[idx]" is a list of lists for the link weights in layer whose o
            transposed_layer_params = list(zip(*layer_params)) ## creating a transpo

            backproped_error = [None] * len(vars_in_next_layer_back)
            for k,varr in enumerate(vars_in_next_layer_back):
                for j,var2 in enumerate(vars_in_layer):
                    backproped_error[k] = sum([self.vals_for_learnable_params[
                        transposed_layer_params[k][i]] *
                        pred_err_backproped_at_layers[back_layer_index][i]
                        for i in range(len(vars_in_layer))])
                    ##      deriv_sigmoid_avg[i] for i in range(len(vars
            pred_err_backproped_at_layers[back_layer_index - 1] = backproped_error
            input_vars_to_layer = self.layer_vars[back_layer_index-1]
            for j,var in enumerate(vars_in_layer):
                layer_params = self.layer_params[back_layer_index][j]

```

```

        for i,param in enumerate(layer_params):
            gradient_of_loss_for_param = input_vals_avg[i] *
            pred_err_backproped_at_layers[back_layer_index][j]
            step = self.mu*self.prev_step[i]+self.learning_rate *
            gradient_of_loss_for_param * deriv_sigmoid_avg[j]
            self.vals_for_learnable_params[param] += step
            self.prev_step[i] = step
        self.step_size = self.mu*self.step_size+self.learning_rate * sum(pred_err_backprop)
        sum(deriv_sigmoid_avg) / len(deriv_sigmoid_avg)
        self.bias[back_layer_index - 1] += self.step_size

def reset_step(self):
    for i in range(len(self.learnable_params)):
        if len(self.prev_step) < len(self.vals_for_learnable_params):
            self.prev_step.append(0)
        else:
            self.prev_step[i] = 0

def run_training_loop_one_neuron_model(self, training_data):
    self.vals_for_learnable_params = {param: random.uniform(0,1)
    for param in self.learnable_params}
    self.bias = [random.uniform(0,1) for _ in range(self.num_layers-1)]
    self.step_size = 0

class DataLoader:
    def __init__(self, training_data, batch_size):
        self.training_data = training_data
        self.batch_size = batch_size
        self.class_0_samples = [(item, 0) for item in self.training_data[0]]
        self.class_1_samples = [(item, 1) for item in self.training_data[1]]
    def __len__(self):
        return len(self.training_data[0]) + len(self.training_data[1])
    def _getitem(self):
        cointoss = random.choice([0,1])
        if cointoss == 0:
            return random.choice(self.class_0_samples)
        else:
            return random.choice(self.class_1_samples)
    def getbatch(self):
        batch_data,batch_labels = [],[]
        maxval = 0.0
        for _ in range(self.batch_size):
            item = self._getitem()
            if np.max(item[0]) > maxval:
                maxval = np.max(item[0])
            batch_data.append(item[0])

```

```

        batch_labels.append(item[1])
        batch_data = [item/maxval for item in batch_data]
        batch = [batch_data, batch_labels]
        return batch

data_loader = DataLoader(training_data, batch_size=self.batch_size)
loss_running_record = []
i = 0
avg_loss_over_literations = 0.0
self.reset_step()

for i in range(self.training_iterations):
    data = data_loader.getbatch()
    data_tuples = data[0]
    class_labels = data[1]
    self.forward_prop_multi_neuron_model(data_tuples)
    predicted_labels_for_batch = self.forw_prop_vals_at_layers[self.num_layers-1]
    y_preds = [item for sublist in predicted_labels_for_batch
               for item in sublist]
    loss = sum([(abs(class_labels[i] - y_preds[i]))**2
               for i in range(len(class_labels))])
    loss_avg = loss / float(len(class_labels))
    avg_loss_over_literations += loss_avg
    if i%(self.display_loss_how_often) == 0:
        avg_loss_over_literations /= self.display_loss_how_often
        loss_running_record.append(avg_loss_over_literations)
        print("[iter=%d] loss = %.4f" % (i+1, avg_loss_over_literations))
        avg_loss_over_literations = 0.0
    y_errors = list(map(operator.sub, class_labels, y_preds))
    y_error_avg = sum(y_errors) / float(len(class_labels))
    self.backprop_and_update_params_multi_neuron_model(y_error_avg, class_labels)

return loss_running_record

cgp = ComputationalGraphPrimerSGDPlus(
    num_layers = 3,
    layers_config = [4,2,1],
    # num of nodes in each layer
    expressions = ["xw=ap*xp+aq*xq+ar*xr+as*xs",
                  "xz=bp*xp+bq*xq+br*xr+bs*xs",
                  "xo=cp*xw+cq*xz"],
    output_vars = ["xo"],

    dataset_size = 5000,
    learning_rate = 1e-3,
    # learning_rate = 5 * 1e-2,

```

```

        training_iterations = 40000,
        batch_size = 8,
        display_loss_how_often = 100,
        debug = True,
        mu = 0
    )

if __name__ == '__main__':
    cgp.parse_multi_layer_expressions()

    training_data = cgp.gen_training_data()

    loss_sgd = cgp.run_training_loop_one_neuron_model( training_data )

    cgp_sgd_plus = ComputationalGraphPrimerSGDPlus(
        num_layers = 3,
        layers_config = [4,2,1],
        # num of nodes in each layer
        expressions = ['xw=ap*xp+aq*xq+ar*xr+as*xs',
            'xz=bp*xp+bq*xq+br*xr+bs*xs',
            'xo=cp*xw+cq*xz'],
        output_vars = ['xo'],

        dataset_size = 5000,
        learning_rate = 1e-3,
        # learning_rate = 5 * 1e-2,
        training_iterations = 40000,
        batch_size = 8,
        display_loss_how_often = 100,
        debug = True,
        mu = 0.99
    )

    cgp_sgd_plus.parse_multi_layer_expressions()

    training_data = cgp_sgd_plus.gen_training_data()

    loss_sgd_plus = cgp_sgd_plus.run_training_loop_one_neuron_model( training_data )

    plt.figure()
    plt.plot(loss_sgd, label = "SGD Training Loss")
    plt.plot(loss_sgd_plus, label = "SGD+ Training Loss")
    plt.xlabel('iterations')
    plt.ylabel('loss')
    plt.title('SGD+ vs SGD Loss Multi Neuron')
    plt.legend()

```



```
plt.show()
```

3.1 Outputs

Figure 1: One Neuron SGD without Optimization with learning rate: 10^{-3} ,batch size:8.




Figure 2: Multi Neuron SGD verified with Torch and learning rate: 10^{-6} ,batch size:8




Figure 3: One Neuron SGD verified with Torch and learning rate: 10^{-3} ,batch size:8




Figure 4: One Neuron SGD comparison with and without optimization and learning rate: 10^{-3} ,batch size:8, $\mu = 0.99$

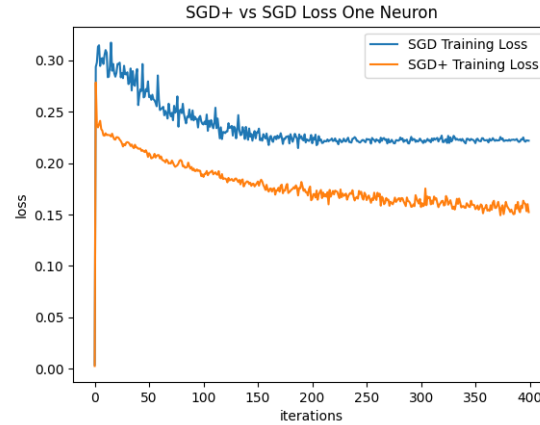



Figure 5: Multi Neuron SGD comparison with and without optimization and learning rate: 10^{-3} , batch size:8, $\mu = 0.99$



Hence optimization using momentum has clearly lead to lower loss and better accuracy.

4 Lessons Learned

- Optimizing SGD using momentum
- Most of the times momentum will work well when slightly above 0.9
- The code to model neural networks from scratch

5 Suggested Enhancements

A lot of the code is still a black box.